

FordAsleep

July 10, 2017

1 Stay Alert! The Ford Challenge

by Scott Josephson

Driving while distracted, fatigued or drowsy may lead to accidents. Activities that divert the driver's attention from the road ahead, such as engaging in a conversation with other passengers in the car, making or receiving phone calls, sending or receiving text messages, eating while driving or events outside the car may cause driver distraction. Fatigue and drowsiness can result from driving long hours or from lack of sleep.

The data for this Kaggle challenge shows the results of a number of "trials", each one representing about 2 minutes of sequential data that are recorded every 100 ms during a driving session on the road or in a driving simulator. The trials are samples from some 100 drivers of both genders, and of different ages and ethnic backgrounds. The files are structured as follows:

The first column is the Trial ID - each period of around 2 minutes of sequential data has a unique trial ID. For instance, the first 1210 observations represent sequential observations every 100ms, and therefore all have the same trial ID. The second column is the observation number - this is a sequentially increasing number within one trial ID. The third column has a value X for each row where

X = 1 if the driver is alert

X = 0 if the driver is not alert

The next 8 columns with headers P1, P2,, P8 represent physiological data;

The next 11 columns with headers E1, E2,, E11 represent environmental data;

The next 11 columns with headers V1, V2,, V11 represent vehicular data;

1.1 Import Libraries

```
In [1]: import numpy as np
import pandas as pd

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report
```

1.2 Get the Data

Read in the `fordtrain.csv` file and set it to a data frame called `ford_train`.

**** Split the data into training set and testing set using `train_test_split` ****

```
In [2]: ford_train = pd.read_csv('fordtrain.csv')
```

Check the head of `ad_data`

```
In [3]: ford_train.head()
```

```
Out[3]:
```

	TrialID	ObsNum	IsAlert	P1	P2	P3	P4	P5	P6	\
0	0	0	0	34.7406	9.84593	1400	42.8571	0.290601	572	
1	0	1	0	34.4215	13.41120	1400	42.8571	0.290601	572	
2	0	2	0	34.3447	15.18520	1400	42.8571	0.290601	576	
3	0	3	0	34.3421	8.84696	1400	42.8571	0.290601	576	
4	0	4	0	34.3322	14.69940	1400	42.8571	0.290601	576	

	P7	...	V2	V3	V4	V5	V6	V7	V8	V9	V10	V11
0	104.895	...	0.175	752	5.99375	0	2005	0	13.4	0	4	14.8004
1	104.895	...	0.455	752	5.99375	0	2007	0	13.4	0	4	14.7729
2	104.167	...	0.280	752	5.99375	0	2011	0	13.4	0	4	14.7736
3	104.167	...	0.070	752	5.99375	0	2015	0	13.4	0	4	14.7667
4	104.167	...	0.175	752	5.99375	0	2017	0	13.4	0	4	14.7757

[5 rows x 33 columns]

```
In [4]: ford_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 604329 entries, 0 to 604328
Data columns (total 33 columns):
TrialID      604329 non-null int64
ObsNum       604329 non-null int64
IsAlert      604329 non-null int64
P1           604329 non-null float64
P2           604329 non-null float64
P3           604329 non-null int64
P4           604329 non-null float64
P5           604329 non-null float64
P6           604329 non-null int64
P7           604329 non-null float64
P8           604329 non-null int64
E1           604329 non-null float64
E2           604329 non-null float64
E3           604329 non-null int64
E4           604329 non-null int64
E5           604329 non-null float64
E6           604329 non-null int64
```

```

E7          604329 non-null int64
E8          604329 non-null int64
E9          604329 non-null int64
E10         604329 non-null int64
E11         604329 non-null float64
V1          604329 non-null float64
V2          604329 non-null float64
V3          604329 non-null int64
V4          604329 non-null float64
V5          604329 non-null int64
V6          604329 non-null int64
V7          604329 non-null int64
V8          604329 non-null float64
V9          604329 non-null int64
V10         604329 non-null int64
V11         604329 non-null float64
dtypes: float64(14), int64(19)
memory usage: 152.2 MB

```

2 Logistic Regression

Now it's time to do a train test split, and train our model!

Choose columns that you want to train on!

```

In [19]: X_train, X_test, y_train, y_test = train_test_split(ford_train.drop('IsAlert',axis=1)
                                                             test_size=0.30,random_state=101)

```

**** Train and fit a logistic regression model on the training set.****

```

In [21]: logmodel = LogisticRegression()

```

```

In [22]: logmodel.fit(X_train, y_train)

```

```

Out[22]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                             intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
                             penalty='l2', random_state=None, solver='liblinear', tol=0.0001,
                             verbose=0, warm_start=False)

```

2.1 Predictions and Evaluations

**** Now predict values for the testing data.****

```

In [23]: predictions = logmodel.predict(X_test)

```

**** Create a classification report for the model.****

```

In [25]: print(classification_report(y_test,predictions))

```

	precision	recall	f1-score	support
0	0.82	0.73	0.77	76334
1	0.82	0.88	0.85	104965
avg / total	0.82	0.82	0.82	181299